

# ExaSAT: Exascale Static Analysis Tool

September 18, 2012

## Abstract

The ExaSAT tool is developed to automatically analyze and compute key performance-related metrics for a given code. The tool is composed of two subtools. The first tool, developed on top of ROSE compiler framework, performs static compiler analysis and focuses on loops and their floating point and memory operations. The collected information from the compiler analysis is stored in an XML format. The XML output is fed to the second tool which generates a spreadsheet containing estimated performance data and a loop dependency graph. The spreadsheet allows users to enter architecture specific parameters such as cache size or memory bandwidth. The spreadsheet will update the estimated time accordingly. We currently implemented the compiler analysis tool for Combustion codes written in Fortran.

## 1 Static Analysis Description

With the help of the ExaSAT tool, we would like to answer questions about the performance of combustion codes on future architectures. Some of the questions we would like our tool to answer are:

- how much data must be moved?
- how sensitive is the code to memory bandwidth?
- what is the fraction of communication in overall running time?
- what is my working set size?
- etc.

### 1.1 Compiler Analysis

We have developed the compiler analysis using the ROSE compiler framework. The compiler analysis tool collects information per function (subroutine or module) and it further collects more detail information for nested loops in a function. The information collected at the function level includes:

The information collected at the loop level includes:




Figure 1: Overview of the ExaSAT tool

## 1.2 XML Description

The results from the first part of the static analysis toolchain are output in an XML format to interface with the other parts of the toolchain. The information in the XML file includes:

- List
- of things
- in the
- XML

Please see the separate XML document detailing the structure of the XML.

## 2 Performance Model

### 2.1 Working Sets and Bandwidth Consumption

*Note: in the following analysis, the cache is assumed to have an ideal LRU eviction policy. Since real architectures do not have ideal LRU policies, the analysis of the cache hit rate and resulting memory traffic is only an approximation of what will be observed in practice.*

Each array may have a different access pattern, so the tool computes working set and bandwidth usage for each array independently given the array's access pattern. The number of planes and pencils in the working set and bandwidth calculations are different because the number of unique items that require space in cache to enable reuse may be different from the number of unique items accessed during a single sweep due to gaps in the access pattern.

The proper working set size that is both necessary and sufficient to enable maximum reuse between pencil or plane iterations equals the span of the pattern plus the maximum gap size across ALL patterns accessed in the loop. For example, a stencil pattern that accesses planes -2, -1, 0, +1, +2, has a working set of 5 planes because there's no gap, but a pattern of -2, -1, +1, +2 requires a working set of 6 planes for reuse between plane sweeps even though it only accesses 4 unique planes during a single sweep. It may seem counter-intuitive




Figure 2: Dependency graph between loops in function `diffterm` in CNS code `advance.f90`.

that accessing fewer planes can increase the working set size, but gaps in the pattern require the cache to “remember” some of the data for a longer period of time without evicting it. This requirement can cause an increase in the working set size.

A similar calculation for pencils is done to enable reuse between pencil sweeps, but since the pattern for pencils is 2D, we consider each unique plane in the pattern when computing the maximum gap and working sets sizes.

## 2.2 Dependency Graph Description

A dependency graph for each function is generated that shows the dependencies between loops (or solvers) in the function. Flow, anti, and output dependencies are considered across all arrays read and written in each loop. An example dependency graph is shown in Figure 2.

## 2.3 Spreadsheet Description

### 2.3.1 Overview

At the top of the spreadsheet, there’s a table of user-modifiable parameters. User-modifiable parameters allow you to change problem parameters including problem size and cache blocking factor in addition to machine architecture parameters like Gflop/s/core, GB/s/chip, and cache/core. The rest of the spreadsheet updates itself to reflect the changes made by the user.

The first section beneath the parameters is a loop-level analysis table listing properties for each loop in the code, including iteration size, block size, flop counts, working set size, estimated bandwidth usage, and estimated execution time for each loop. Totals for the whole program are also included.

The second section breaks out the analysis for each of the read-only arrays in each loop because they have ghost cell accesses. The numbers computed in

this section are then used in the total working set and bandwidth calculations in the first section.

### 2.3.2 Parameters

The parameters at the top allow the user to explore different settings:

Changing the blocking factor may decrease execution time if it enables greater reuse of cached data. On the other hand, increasing the blocking factor by too great an amount could hurt performance, as there is increased memory bandwidth consumed when pulling in the ghost cells for more blocks.

Changing the cache available per core may affect what type of reuse occurs in each loop of the program. The spreadsheet compares the available machine cache size with the working set sizes required for reuse between pencil iterations and plane iterations.

The cost of reads (R), read-writes (RW), and writes (W) affect the bandwidth consumed by each of the memory operations. For example, in a cache-bypass setting, the write bandwidth could be half that of the read-write bandwidth, whereas if no cache-bypass is used, it could be equal to the read-write bandwidth.

The last two parameters allows two optimizations to be made to how the cache is managed. If Streaming Writes is TRUE, it is assumed that data from read-write and write-only arrays are streamed into registers and do not pollute the cache. If the additional optimization of using NTA cache hints is used, it is assumed that streaming reads (reads with no reuse) will not pollute the cache either. These optimizations affect the computed working set sizes required for different types of reuse (reuse between plane sweeps and reuse between pencil sweeps).

### 2.3.3 Loop Analysis Table

This table lists the following properties for each loop in the code:

- Function name and line number of the loop
- Number of sweeps (e.g. if it is run for each component in an array)
- Total and block iteration space
- Number of floating point operations per iteration (adds, multiplies, and specials)
- Number of arrays for read-write (RW) and write-only (W) access
- Number of planes and pencils required in the working set to enable reuse between sweeps
- Four pairs of columns, each pair listing 1) the working set in MB to enable reuse between plane sweeps and 2) the working set in MB to enable reuse between pencil sweeps. The first pair is for a naïve memory access pattern.

The second pair lists working sets required if streaming writes are utilized. The third pair lists the optimal working set where only data that is reused resides in cache. The fourth pair lists the actual working sets based on the user's selections in the parameter table section. If NTA Hints is TRUE, then the fourth pair of columns points to the third pair (reuse only). If NTA Hints is FALSE, but Streaming Writes is TRUE, then the fourth pair points to the second pair (streaming writes). If both are FALSE, then the fourth pair points to the first pair (naive).

- Four columns showing the memory bandwidth consumed under different scenarios. The first column shows the amount of data transferred if there is reuse between plane sweeps of the cache block. This corresponds to the compulsory traffic for the cache block. The second column shows amount transferred if there is reuse between pencil sweeps, but not plane sweeps. The third column shows the amount transferred if there is no reuse between pencil sweeps. The fourth column shows the actual amount transferred given the size of the cache specified in the parameters table at the top of the spreadsheet and the working set sizes required for different types of reuse given in the "Working set (actual)" columns of this table.
- Gflops performed per iteration and the arithmetic intensity of the computation given the Gflops performed and the Gbytes transferred
- The estimated execution times if the program is CPU limited (CPU), or memory bandwidth limited (DRAM). The final estimate (CPU and DRAM) is the maximum between these two values.